

March 1, 2023

```
[ ]: # ! pip install contractions
# ! pip install scikit-learn
# ! pip install pandas
# ! pip install numpy
# ! pip install nltk
# ! pip install gensim
# ! pip3 install torch torchvision torchaudio --extra-index-url https://
↳download.pytorch.org/whl/cu116
# # Dataset: https://s3.amazonaws.com/amazon-reviews-pds/tsv/
# # amazon_reviews_us_Beauty_v1_00.tsv.gz
```

```
[ ]: import pandas as pd
import numpy as np
import re
from gensim.models import Word2Vec
import gensim.downloader as api
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report
from sklearn.linear_model import Perceptron
from sklearn.svm import LinearSVC
import torch
import torch.nn as nn
import torch.nn.functional as F
from torch.utils.data import Dataset
from torch.utils.data import DataLoader

import nltk
nltk.download('punkt')
nltk.download('wordnet')
from nltk.corpus import wordnet
from sklearn.feature_extraction.text import TfidfVectorizer
```

```
c:\Users\karav\Desktop\Applied NLP\HW3\venv\lib\site-packages\tqdm\auto.py:22:
TqdmWarning: IProgress not found. Please update jupyter and ipywidgets. See
https://ipywidgets.readthedocs.io/en/stable/user_install.html
```

```
from .autonotebook import tqdm as notebook_tqdm
[nltk_data] Downloading package punkt to
[nltk_data] C:\Users\karav\AppData\Roaming\nltk_data...
```

```
[nltk_data] Package punkt is already up-to-date!
[nltk_data] Downloading package wordnet to
[nltk_data] C:\Users\karav\AppData\Roaming\nltk_data...
[nltk_data] Package wordnet is already up-to-date!
```

```
[ ]: # data = pd.read_csv('amazon_reviews_us_Beauty_v1_00.tsv', sep='\t',
    ↪encoding='utf-8', on_bad_lines='skip')
```

```
[ ]: # data['star_rating'] = data['star_rating'].astype('int', errors='ignore')

# df = data[['review_body', 'star_rating']]
# df = df.dropna(thresh=2)
```

Form three classes of 20000 reviews randomly

```
[ ]: # class_a = df.loc[df['star_rating'].isin([1,2]).sample(n=20000,
    ↪random_state=4)
# class_b = df.loc[df['star_rating'].isin([3]).sample(n=20000, random_state=7)
# class_c = df.loc[df['star_rating'].isin([4,5]).sample(n=20000,
    ↪random_state=21)

# df_sampled = pd.concat([class_a, class_b, class_c])
# df_sampled['star_rating'].value_counts()
```

```
[ ]: # df_sampled['star_rating'] = df_sampled['star_rating'].replace([1,2],0)
# df_sampled['star_rating'] = df_sampled['star_rating'].replace([3],1)
# df_sampled['star_rating'] = df_sampled['star_rating'].replace([4,5],2)

# df_sampled.to_csv("data.tsv", sep='\t')
```

So far, I have read the dataset from the tsv file. I converted the data type of star_rating field to int in order to enforce uniformity as there were some float and date values in them. After dropping null values, sampled 20000 reviews for each class label and merged them into a single dataframe and then stored it for easier read access.

0.1 Task 2

```
[ ]: w2v = api.load('word2vec-google-news-300')
```

```
[ ]: # Functions to clean dataset by removing html, urls, punctuation and multiple
    ↪spaces

def remove_HTML(s):
    return re.sub(r'<.*?>', ' ', s)

def remove_URL(s):
    return re.sub(r'https?:\\\/\\\/.*\\\/w*', ' ', s)
```

```
def remove_nonalphabets(s):
    return re.sub(r'[^a-zA-Z]', ' ', s)

def remove_multispace(s):
    return re.sub(r'\s+', ' ', s)
```

```
[ ]: data = pd.read_csv('data.tsv', sep='\t', usecols=['review_body',
↳ 'star_rating'], dtype={'review_body' : str, 'star_rating': int})

data['review_body'] = data['review_body'].apply(lambda x:remove_HTML(x))
data['review_body'] = data['review_body'].apply(lambda x:remove_URL(x))
data['review_body'] = data['review_body'].apply(lambda x:remove_nonalphabets(x))
data['review_body'] = data['review_body'].apply(lambda x:remove_multispace(x))
tfidfdata = data.copy(deep=True)
data['review_body'] = data['review_body'].apply(lambda x: x.split())
tfidfdata['review_body'] = tfidfdata['review_body'].apply(lambda x: x.lower().
↳ split())
tfidfdata['review_body'] = tfidfdata['review_body'].apply(lambda x: ' '.join(x))

w2vdata = list(data['review_body'])

[ ]: model = Word2Vec(w2vdata, min_count=9, vector_size=300, window=13, workers=6)
```

0.1.1 2(a)

```
[ ]: #Testing examples on Google's word2vec

test_vec = w2v['China']-w2v['India']+w2v['Indian']
print(w2v.most_similar(positive=[test_vec],topn=3))
print(w2v.similarity('excellent','outstanding'))
print(w2v.similarity('excellent','poor'))
print(w2v.similarity('beautiful','horrible'))
print(w2v.doesnt_match(['excellent','outstanding','poor']))
print(w2v.doesnt_match(['excellent','bad','poor']))

[('Chinese', 0.8240145444869995), ('China', 0.6561532616615295), ('Indian',
0.643246054649353)]
0.55674857
0.37769592
0.38830176
poor
excellent
```

0.1.2 2(b)

```
[ ]: #Testing examples on trained word2vec

test_vec = model.wv['China']-model.wv['India']+model.wv['Indian']
print(model.wv.most_similar(positive=[test_vec],topn=3))
print(model.wv.similarity('excellent','outstanding'))
print(model.wv.similarity('excellent','poor'))
print(model.wv.similarity('beautiful','horrible'))
print(model.wv.doesnt_match(['excellent','outstanding','poor']))
print(model.wv.doesnt_match(['excellent','bad','poor']))

[('China', 0.9744254350662231), ('china', 0.7366096377372742), ('USA',
0.7354326248168945)]
0.73620814
0.5710197
0.40271857
poor
bad
```

```
[ ]: del model
```

The vectors generated by our dataset do not seem to capture the relationships between words accurately. Although it seems to place words that are relevant to our context in proximity in vector space, the relationships are not encoded properly. It is not capable of performing the vector algebra or picking odd one out as well as pretrained word2vec.

0.2 Task 3

```
[ ]: Vectorizer = TfidfVectorizer()
tfidf = Vectorizer.fit_transform(tfidfdata['review_body'])

TF_train, TF_test, tf_train, tf_test = train_test_split(tfidf,
                                                         tfidfdata['star_rating'],
                                                         stratify=tfidfdata['star_rating'],
                                                         test_size=0.2, random_state=1)

p = Perceptron(random_state=7)
p.fit(TF_train, tf_train)

print("-----TF-IDF Features-----")
print('Perceptron')
print(classification_report(tf_test, p.predict(TF_test)))

print('-----')
print('SVM')
s = LinearSVC(random_state=7, tol= 1e-5)
s.fit(TF_train, tf_train)
print(classification_report(tf_test, s.predict(TF_test)))
```

-----TF-IDF Features-----

Perceptron

	precision	recall	f1-score	support
0	0.62	0.69	0.65	4000
1	0.56	0.51	0.53	4000
2	0.71	0.70	0.70	4000
accuracy			0.63	12000
macro avg	0.63	0.63	0.63	12000
weighted avg	0.63	0.63	0.63	12000

SVM

	precision	recall	f1-score	support
0	0.71	0.71	0.71	4000
1	0.61	0.59	0.60	4000
2	0.76	0.79	0.77	4000
accuracy			0.70	12000
macro avg	0.69	0.70	0.69	12000
weighted avg	0.69	0.70	0.69	12000

```
[ ]: del tfidfdata, Vectorizer, tfidf, TF_train, tf_train, TF_test, tf_test
```

```
[ ]: def transform_w2v(data):
    w2vfeats = []
    for sentence in data:
        word_vecs = [w2v[word] for word in sentence if word in w2v]
        if len(word_vecs):
            sent_vec = np.mean(word_vecs, axis=0)
        else:
            sent_vec = np.zeros(300)
        w2vfeats.append(sent_vec)
    return np.array(w2vfeats)

w2vfeats_np = transform_w2v(w2vdata)
w2vlabels_np = np.array(list(data['star_rating'].astype('int')))
```

I have averaged all the word vectors for each review. For reviews that have no words in the w2v vocabulary, the feature vector is a list of zeros.

```
[ ]: X_train, X_test, y_train, y_test = train_test_split(w2vfeats_np,
                                                         w2vlabels_np,
                                                         stratify=w2vlabels_np,
                                                         test_size=0.2, random_state=1)
```

```
[ ]: print("-----Word2Vec Features-----")
print('Perceptron')
p = Perceptron(random_state=7)
p.fit(X_train, y_train)
print(classification_report(y_test, p.predict(X_test)))

print("-----")
print('SVM')
s = LinearSVC(random_state=7, tol= 1e-5)
s.fit(X_train, y_train)
print(classification_report(y_test, s.predict(X_test)))

del p,s
```

-----Word2Vec Features-----

Perceptron

	precision	recall	f1-score	support
0	0.79	0.19	0.30	4000
1	0.39	0.93	0.55	4000
2	0.87	0.33	0.48	4000
accuracy			0.48	12000
macro avg	0.68	0.48	0.44	12000
weighted avg	0.68	0.48	0.44	12000

SVM

	precision	recall	f1-score	support
0	0.65	0.68	0.67	4000
1	0.59	0.55	0.57	4000
2	0.71	0.72	0.72	4000
accuracy			0.65	12000
macro avg	0.65	0.65	0.65	12000
weighted avg	0.65	0.65	0.65	12000

TF-IDF features -

Perceptron : 63% accuracy | SVM : 70% accuracy

Word2Vec features -

Perceptron : 48% accuracy | SVM - 65% accuracy

The models were able to classify better based on TF-IDF features than the word2vec features. This could be due to few reasons - TF-IDF is a statistical method that is intended to improve metrics such as precision and recall. Word2Vec is aimed at capturing semantic relationships. Google's pretrained word2vec was trained on a wide context that may not have seen all these emotions to weigh them well in vector space.

```
[ ]: # if torch.cuda.is_available():
#     device = torch.device("cuda")
#     print(device)
# else:
device = torch.device("cpu")
```

0.3 Task 4

```
[ ]: class TrainDataset(Dataset):
    def __init__(self, reviews, labels, transform=None):
        self.reviews = torch.from_numpy(reviews).float()
        self.labels = torch.from_numpy(labels).long()
        self.transform = transform

    def __len__(self):
        return len(self.reviews)

    def __getitem__(self, idx):
        review = self.reviews[idx]
        label = self.labels[idx]
        if self.transform:
            review = self.transform(review)
```

```
return review, label
```

```
[ ]: params = {
    'batch_size': 64,
    'shuffle': True,
    'num_workers': 0
}

train_data = TrainDataset(X_train, y_train)
train_loader = DataLoader(train_data, **params)

valid_data = TrainDataset(X_test, y_test)
valid_loader = DataLoader(valid_data, **params)
```

Created a custom Dataset class inheriting from torch's Dataset and used a dataloader for batching

0.3.1 4(a)

```
[ ]: class MLP1(nn.Module):
    def __init__(self):
        super(MLP1, self).__init__()
        self.fc1 = nn.Linear(300, 100)
        self.fc2 = nn.Linear(100, 10)
        self.fc3 = nn.Linear(10, 3)
        self.dropout = nn.Dropout(0.2)

    def forward(self, x):
        x = F.relu(self.fc1(x))
        x = self.dropout(x)
        x = F.relu(self.fc2(x))
        x = self.dropout(x)
        x = F.log_softmax(self.fc3(x), dim=1)
        # x = self.fc3(x)
        return x

mlp1 = MLP1()
mlp1 = mlp1.to(device)
loss_fn = nn.CrossEntropyLoss()
optimizer = torch.optim.AdamW(mlp1.parameters(), lr=0.01)

[ ]: max_epochs = 30
valid_loss_min = np.Inf

for epoch in range(max_epochs):
    train_loss = 0
    valid_loss = 0
    mlp1.train()
```



```

train_correct=0
for batch, labels in train_loader:
    batch, labels = batch.to(device), labels.to(device)
    optimizer.zero_grad()
    output = mlp1(batch)
    loss = loss_fn(output, labels)
    predicted = torch.argmax(output, dim=1)
    loss.backward()
    optimizer.step()
    train_loss+= loss.item()
    train_correct+= (predicted==labels).sum().item()

mlp1.eval()
valid_correct=0
with torch.no_grad():
    for batch, labels in valid_loader:
        batch, labels = batch.to(device), labels.to(device)
        output = mlp1(batch)
        loss = loss_fn(output, labels)
        predicted = torch.argmax(output, dim=1)
        valid_loss+= loss.item()
        valid_correct+=(predicted==labels).sum().item()

train_loss = train_loss/len(train_loader)
valid_loss = valid_loss/len(valid_loader)
train_acc = 100 * train_correct / len(train_data)
valid_acc = 100 * valid_correct / len(valid_data)

print('Epoch: {} \tTraining Loss: {:.6f} \tValidation Loss: {:.6f}\n
↪\tTraining Accuracy: {:.6f} \t Validation Accuracy: {:.6f}'.format(
    epoch+1,
    train_loss,
    valid_loss,
    train_acc,
    valid_acc
))

if valid_loss <= valid_loss_min:
    print('Validation loss decreased ({:.6f} --> {:.6f}). Saving model ...
↪'.format(
        valid_loss_min,
        valid_loss))
    torch.save(mlp1.state_dict(), 'mlp1.pt')
    valid_loss_min = valid_loss

del mlp1

```

Epoch: 1 Training Loss: 0.882419 Validation Loss: 0.805861
Training Accuracy: 58.970833 Validation Accuracy: 64.666667
Validation loss decreased (inf --> 0.805861). Saving model ...

Epoch: 2 Training Loss: 0.838510 Validation Loss: 0.840016
Training Accuracy: 62.397917 Validation Accuracy: 62.116667

Epoch: 3 Training Loss: 0.823994 Validation Loss: 0.791876
Training Accuracy: 62.995833 Validation Accuracy: 64.216667
Validation loss decreased (0.805861 --> 0.791876). Saving model ...

Epoch: 4 Training Loss: 0.820845 Validation Loss: 0.786919
Training Accuracy: 63.077083 Validation Accuracy: 64.991667
Validation loss decreased (0.791876 --> 0.786919). Saving model ...

Epoch: 5 Training Loss: 0.811417 Validation Loss: 0.802333
Training Accuracy: 63.672917 Validation Accuracy: 65.258333

Epoch: 6 Training Loss: 0.803838 Validation Loss: 0.788787
Training Accuracy: 64.104167 Validation Accuracy: 64.850000

Epoch: 7 Training Loss: 0.800308 Validation Loss: 0.787960
Training Accuracy: 64.539583 Validation Accuracy: 63.950000

Epoch: 8 Training Loss: 0.799988 Validation Loss: 0.773323
Training Accuracy: 64.083333 Validation Accuracy: 65.500000
Validation loss decreased (0.786919 --> 0.773323). Saving model ...

Epoch: 9 Training Loss: 0.795391 Validation Loss: 0.783199
Training Accuracy: 64.558333 Validation Accuracy: 65.475000

Epoch: 10 Training Loss: 0.793381 Validation Loss: 0.767247
Training Accuracy: 64.854167 Validation Accuracy: 65.966667
Validation loss decreased (0.773323 --> 0.767247). Saving model ...

Epoch: 11 Training Loss: 0.792889 Validation Loss: 0.769620
Training Accuracy: 64.683333 Validation Accuracy: 65.750000

Epoch: 12 Training Loss: 0.789712 Validation Loss: 0.780824
Training Accuracy: 64.731250 Validation Accuracy: 64.975000

Epoch: 13 Training Loss: 0.788710 Validation Loss: 0.782966
Training Accuracy: 64.991667 Validation Accuracy: 64.800000

Epoch: 14 Training Loss: 0.787117 Validation Loss: 0.771606
Training Accuracy: 64.793750 Validation Accuracy: 65.833333

Epoch: 15 Training Loss: 0.787902 Validation Loss: 0.774148
Training Accuracy: 65.154167 Validation Accuracy: 65.391667

Epoch: 16 Training Loss: 0.781290 Validation Loss: 0.775664
Training Accuracy: 65.266667 Validation Accuracy: 65.608333

Epoch: 17 Training Loss: 0.779198 Validation Loss: 0.769034
Training Accuracy: 65.466667 Validation Accuracy: 66.341667

Epoch: 18 Training Loss: 0.780767 Validation Loss: 0.762907
Training Accuracy: 65.554167 Validation Accuracy: 66.291667
Validation loss decreased (0.767247 --> 0.762907). Saving model ...

Epoch: 19 Training Loss: 0.780964 Validation Loss: 0.763635
Training Accuracy: 65.177083 Validation Accuracy: 65.883333

Epoch: 20 Training Loss: 0.780560 Validation Loss: 0.796080
Training Accuracy: 65.337500 Validation Accuracy: 63.791667

Epoch: 21 Training Loss: 0.776748 Validation Loss: 0.764099
Training Accuracy: 65.718750 Validation Accuracy: 66.158333

```

Epoch: 22      Training Loss: 0.780137      Validation Loss: 0.767298
Training Accuracy: 65.433333      Validation Accuracy: 66.300000
Epoch: 23      Training Loss: 0.771557      Validation Loss: 0.765212
Training Accuracy: 65.685417      Validation Accuracy: 65.925000
Epoch: 24      Training Loss: 0.774043      Validation Loss: 0.765254
Training Accuracy: 65.566667      Validation Accuracy: 66.341667
Epoch: 25      Training Loss: 0.772692      Validation Loss: 0.799089
Training Accuracy: 65.539583      Validation Accuracy: 65.008333
Epoch: 26      Training Loss: 0.773041      Validation Loss: 0.766370
Training Accuracy: 65.804167      Validation Accuracy: 66.200000
Epoch: 27      Training Loss: 0.770929      Validation Loss: 0.765811
Training Accuracy: 65.929167      Validation Accuracy: 66.216667
Epoch: 28      Training Loss: 0.773404      Validation Loss: 0.773933
Training Accuracy: 65.641667      Validation Accuracy: 65.400000
Epoch: 29      Training Loss: 0.773131      Validation Loss: 0.784065
Training Accuracy: 65.754167      Validation Accuracy: 65.333333
Epoch: 30      Training Loss: 0.773960      Validation Loss: 0.760797
Training Accuracy: 65.477083      Validation Accuracy: 66.108333
Validation loss decreased (0.762907 --> 0.760797). Saving model ...

```

```
[ ]: #model.load_state_dict(torch.load('mlp1.pt')) to load best valid_loss model
```

0.3.2 4(b)

```
[ ]: def transform_w2v_10words(data):
    w2vfeats = []
    for sentence in data:
        word_vecs = [w2v[word] for word in sentence if word in w2v]
        sent_vec = np.concatenate([word_vecs[i] if i < len(word_vecs) else np.
        ↪zeros(300) for i in range(10)], axis=0)
        w2vfeats.append(sent_vec)
    return np.array(w2vfeats)

w2vfeats_np = transform_w2v_10words(w2vdata)
```

Concatenated only first 10 words to generate a vector for reviews.

```
[ ]: X_train, X_test, y_train, y_test = train_test_split(w2vfeats_np,
    w2vlabels_np,
    stratify=w2vlabels_np,
    test_size=0.2, random_state=1)
```

```
[ ]: params = {
    'batch_size': 64,
    'shuffle': True,
    'num_workers': 0
}
```

```

train_data = TrainDataset(X_train, y_train)
train_loader = DataLoader(train_data, **params)

valid_data = TrainDataset(X_test, y_test)
valid_loader = DataLoader(valid_data, **params)

```

```

[ ]: class MLP2(nn.Module):
    def __init__(self):
        super(MLP2, self).__init__()
        self.fc1 = nn.Linear(3000, 100)
        self.fc2 = nn.Linear(100, 10)
        self.fc3 = nn.Linear(10, 3)
        self.dropout = nn.Dropout(0.3)

    def forward(self, x):
        x = F.relu(self.fc1(x))
        x = self.dropout(x)
        x = F.relu(self.fc2(x))
        x = self.dropout(x)
        x = F.log_softmax(self.fc3(x), dim=1)
        return x

mlp2 = MLP2()
mlp2 = mlp2.to(device)
loss_fn = nn.CrossEntropyLoss()
optimizer = torch.optim.AdamW(mlp2.parameters(), lr=0.005)

```

```

[ ]: max_epochs = 30
valid_loss_min = np.Inf

for epoch in range(max_epochs):
    train_loss = 0
    valid_loss = 0
    mlp2.train()
    train_correct=0
    for batch, labels in train_loader:
        batch, labels = batch.to(device), labels.to(device)
        optimizer.zero_grad()
        output = mlp2(batch)
        loss = loss_fn(output, labels)
        predicted = torch.argmax(output, dim=1)
        loss.backward()
        optimizer.step()
        train_loss+= loss.item()
        train_correct+= (predicted==labels).sum().item()

    mlp2.eval()

```

```

valid_correct=0
with torch.no_grad():
    for batch, labels in valid_loader:
        batch, labels = batch.to(device), labels.to(device)
        output = mlp2(batch)
        loss = loss_fn(output, labels)
        predicted = torch.argmax(output, dim=1)
        valid_loss+= loss.item()
        valid_correct+=(predicted==labels).sum().item()

train_loss = train_loss/len(train_loader)
valid_loss = valid_loss/len(valid_loader)
train_acc = 100 * train_correct / len(train_data)
valid_acc = 100 * valid_correct / len(valid_data)

print('Epoch: {} \tTraining Loss: {:.6f} \tValidation Loss: {:.6f}\n
↪\tTraining Accuracy: {:.6f} \t Validation Accuracy: {:.6f}'.format(
    epoch+1,
    train_loss,
    valid_loss,
    train_acc,
    valid_acc
))

if valid_loss <= valid_loss_min:
    print('Validation loss decreased ({:.6f} --> {:.6f}). Saving model ...
↪'.format(
        valid_loss_min,
        valid_loss))
    torch.save(mlp2.state_dict(), 'mlp2.pt')
    valid_loss_min = valid_loss

del mlp2

```

```

Epoch: 1      Training Loss: 0.976678      Validation Loss: 0.924451
Training Accuracy: 50.704167      Validation Accuracy: 54.841667
Validation loss decreased (inf --> 0.924451). Saving model ...
Epoch: 2      Training Loss: 0.917402      Validation Loss: 0.902421
Training Accuracy: 56.052083      Validation Accuracy: 56.175000
Validation loss decreased (0.924451 --> 0.902421). Saving model ...
Epoch: 3      Training Loss: 0.884060      Validation Loss: 0.900157
Training Accuracy: 57.893750      Validation Accuracy: 56.175000
Validation loss decreased (0.902421 --> 0.900157). Saving model ...
Epoch: 4      Training Loss: 0.850590      Validation Loss: 0.909461
Training Accuracy: 59.993750      Validation Accuracy: 56.275000
Epoch: 5      Training Loss: 0.820458      Validation Loss: 0.920390
Training Accuracy: 61.864583      Validation Accuracy: 56.175000

```

Epoch: 6	Training Loss: 0.790396	Validation Loss: 0.934739
Training Accuracy: 63.789583	Validation Accuracy: 54.891667	
Epoch: 7	Training Loss: 0.762101	Validation Loss: 0.934167
Training Accuracy: 65.118750	Validation Accuracy: 55.950000	
Epoch: 8	Training Loss: 0.743555	Validation Loss: 0.948798
Training Accuracy: 66.289583	Validation Accuracy: 56.308333	
Epoch: 9	Training Loss: 0.715384	Validation Loss: 0.999014
Training Accuracy: 67.735417	Validation Accuracy: 55.400000	
Epoch: 10	Training Loss: 0.692863	Validation Loss: 0.995963
Training Accuracy: 69.102083	Validation Accuracy: 54.825000	
Epoch: 11	Training Loss: 0.673906	Validation Loss: 0.999969
Training Accuracy: 69.691667	Validation Accuracy: 55.408333	
Epoch: 12	Training Loss: 0.662720	Validation Loss: 1.002574
Training Accuracy: 70.443750	Validation Accuracy: 55.166667	
Epoch: 13	Training Loss: 0.643923	Validation Loss: 1.020750
Training Accuracy: 71.410417	Validation Accuracy: 54.875000	
Epoch: 14	Training Loss: 0.632360	Validation Loss: 1.061070
Training Accuracy: 72.102083	Validation Accuracy: 55.141667	
Epoch: 15	Training Loss: 0.622716	Validation Loss: 1.056160
Training Accuracy: 72.402083	Validation Accuracy: 55.425000	
Epoch: 16	Training Loss: 0.607511	Validation Loss: 1.085207
Training Accuracy: 73.291667	Validation Accuracy: 55.341667	
Epoch: 17	Training Loss: 0.598487	Validation Loss: 1.114869
Training Accuracy: 73.583333	Validation Accuracy: 55.158333	
Epoch: 18	Training Loss: 0.591536	Validation Loss: 1.111149
Training Accuracy: 73.964583	Validation Accuracy: 54.691667	
Epoch: 19	Training Loss: 0.581423	Validation Loss: 1.144386
Training Accuracy: 74.447917	Validation Accuracy: 54.616667	
Epoch: 20	Training Loss: 0.579383	Validation Loss: 1.146685
Training Accuracy: 74.779167	Validation Accuracy: 54.716667	
Epoch: 21	Training Loss: 0.569262	Validation Loss: 1.178676
Training Accuracy: 75.316667	Validation Accuracy: 55.358333	
Epoch: 22	Training Loss: 0.561153	Validation Loss: 1.165771
Training Accuracy: 75.602083	Validation Accuracy: 55.191667	
Epoch: 23	Training Loss: 0.553000	Validation Loss: 1.181841
Training Accuracy: 75.793750	Validation Accuracy: 53.525000	
Epoch: 24	Training Loss: 0.550625	Validation Loss: 1.216485
Training Accuracy: 76.018750	Validation Accuracy: 54.425000	
Epoch: 25	Training Loss: 0.541611	Validation Loss: 1.220801
Training Accuracy: 76.566667	Validation Accuracy: 54.675000	
Epoch: 26	Training Loss: 0.539132	Validation Loss: 1.136490
Training Accuracy: 76.789583	Validation Accuracy: 54.591667	
Epoch: 27	Training Loss: 0.530489	Validation Loss: 1.258120
Training Accuracy: 77.212500	Validation Accuracy: 54.616667	
Epoch: 28	Training Loss: 0.533744	Validation Loss: 1.198413
Training Accuracy: 76.943750	Validation Accuracy: 54.591667	
Epoch: 29	Training Loss: 0.526562	Validation Loss: 1.253449
Training Accuracy: 77.431250	Validation Accuracy: 54.316667	

Epoch: 30 Training Loss: 0.524036 Validation Loss: 1.316890
Training Accuracy: 77.437500 Validation Accuracy: 55.033333

Observed that in the case of 4(b), the training loss decreases while validation loss increases. Experimented with dropout and lower learning rates but this phenomenon seems to occur regardless. The model in 4(a) performs significantly better than a single perceptron. Adding complexities to the model with non-linearity and layers has improved the performance from 48% to 66.1%. The disparity between 4(a) and 4(b) is possibly because 4(a) tries to capture the entirety of a review while still maintaining lower dimensions whereas 4(b) concatenates the first 10 words, which may not contain all necessary information to classify, and also increases the dimensions of features.

0.4 Task 5

```
[ ]: def transform_rnn(data):  
    w2vfeats = []  
    for sentence in data:  
        word_vecs = [w2v[word] for word in sentence if word in w2v]  
        sent_vec = [word_vecs[i] if i < len(word_vecs) else np.zeros(300) for i  
↪in range(20)]  
        w2vfeats.append(sent_vec)  
    return np.array(w2vfeats)  
  
w2vfeats_np = transform_rnn(w2vdata)
```

Created word2vec features for review by appending 20 word vectors individually without performing any mathematical operations. If there are less than 20 words, the feature is padded with zeros; If there are more than 20, review is truncated at 20 words. This is the input required for RNNs.

```
[ ]: X_train, X_test, y_train, y_test = train_test_split(w2vfeats_np,  
                                                         w2vlabels_np,  
                                                         stratify=w2vlabels_np,  
                                                         test_size=0.2, random_state=1)
```

```
[ ]: params = {  
    'batch_size': 64,  
    'shuffle': True,  
    'num_workers': 0  
}  
  
train_data = TrainDataset(X_train, y_train)  
train_loader = DataLoader(train_data, **params)  
  
valid_data = TrainDataset(X_test, y_test)  
valid_loader = DataLoader(valid_data, **params)
```

0.4.1 5(a)

```
[ ]: class RNN(nn.Module):
    def __init__(self, input_size, hidden_size, output_size):
        super(RNN, self).__init__()
        self.hidden_size = hidden_size
        self.rnn = nn.RNN(input_size, hidden_size, 1, batch_first=True,
        ↪nonlinearity='relu')
        self.fc = nn.Linear(hidden_size, output_size)

    def forward(self, x):
        batch_size = x.size(0)
        hidden = torch.zeros(1, batch_size, self.hidden_size).to(device)
        out, hidden = self.rnn(x, hidden)
        out = F.log_softmax(self.fc(out[:,-1,:]), dim=1) #Last output alone to
        ↪calculate loss
        return out

rnn = RNN(300, 20, 3)
rnn = rnn.to(device)
loss_fn = nn.CrossEntropyLoss()
optimizer = torch.optim.AdamW(rnn.parameters(), lr=0.005)
```

```
[ ]: max_epochs = 30
valid_loss_min = np.Inf

for epoch in range(max_epochs):
    train_loss = 0
    valid_loss = 0
    rnn.train()

    train_correct=0
    for batch, labels in train_loader:
        batch, labels = batch.to(device), labels.to(device)
        optimizer.zero_grad()
        output = rnn(batch)
        # print(output.shape, '/', labels.shape)
        loss = loss_fn(output, labels)
        loss.backward()
        optimizer.step()
        train_loss+= loss.item()
        train_correct+= (torch.argmax(output, dim=1)==labels).sum().item()

    rnn.eval()
    valid_correct=0
    with torch.no_grad():
        for batch, labels in valid_loader:
```



```

        batch, labels = batch.to(device), labels.to(device)
        output = rnn(batch)
        loss=loss_fn(output, labels)
        valid_loss+= loss.item()
        valid_correct+= (torch.argmax(output, dim=1)==labels).sum().item()

train_loss = train_loss/len(train_loader)
valid_loss = valid_loss/len(valid_loader)
train_acc = 100 * train_correct / len(train_data)
valid_acc = 100 * valid_correct / len(valid_data)

print('Epoch: {} \tTraining Loss: {:.6f} \tValidation Loss: {:.6f}\n
↪\tTraining Accuracy: {:.6f} \t Validation Accuracy: {:.6f}'.format(
    epoch+1,
    train_loss,
    valid_loss,
    train_acc,
    valid_acc
))

if valid_loss <= valid_loss_min:
    print('Validation loss decreased ({:.6f} --> {:.6f}). Saving model ...
↪'.format(
        valid_loss_min,
        valid_loss))
    torch.save(rnn.state_dict(), 'rnn.pt')
    valid_loss_min = valid_loss

```

```

Epoch: 1      Training Loss: 0.973483      Validation Loss: 0.924852
Training Accuracy: 50.408333      Validation Accuracy: 55.008333
Validation loss decreased (inf --> 0.924852). Saving model ...
Epoch: 2      Training Loss: 0.884785      Validation Loss: 0.866593
Training Accuracy: 57.091667      Validation Accuracy: 57.658333
Validation loss decreased (0.924852 --> 0.866593). Saving model ...
Epoch: 3      Training Loss: 0.859548      Validation Loss: 0.863384
Training Accuracy: 59.220833      Validation Accuracy: 58.991667
Validation loss decreased (0.866593 --> 0.863384). Saving model ...
Epoch: 4      Training Loss: 0.841908      Validation Loss: 0.880996
Training Accuracy: 60.531250      Validation Accuracy: 58.741667
Epoch: 5      Training Loss: 0.828599      Validation Loss: 0.850196
Training Accuracy: 61.337500      Validation Accuracy: 60.008333
Validation loss decreased (0.863384 --> 0.850196). Saving model ...
Epoch: 6      Training Loss: 0.822405      Validation Loss: 0.841068
Training Accuracy: 61.489583      Validation Accuracy: 60.483333
Validation loss decreased (0.850196 --> 0.841068). Saving model ...
Epoch: 7      Training Loss: 0.811936      Validation Loss: 0.874489
Training Accuracy: 62.456250      Validation Accuracy: 57.950000

```

Epoch: 8 Training Loss: 0.805039 Validation Loss: 0.833688
 Training Accuracy: 63.006250 Validation Accuracy: 60.850000
 Validation loss decreased (0.841068 --> 0.833688). Saving model ...
 Epoch: 9 Training Loss: 0.802850 Validation Loss: 0.843161
 Training Accuracy: 62.750000 Validation Accuracy: 60.425000
 Epoch: 10 Training Loss: 0.797369 Validation Loss: 0.834855
 Training Accuracy: 63.375000 Validation Accuracy: 60.158333
 Epoch: 11 Training Loss: 0.795914 Validation Loss: 0.861628
 Training Accuracy: 63.385417 Validation Accuracy: 60.258333
 Epoch: 12 Training Loss: 0.799824 Validation Loss: 0.824695
 Training Accuracy: 63.541667 Validation Accuracy: 61.400000
 Validation loss decreased (0.833688 --> 0.824695). Saving model ...
 Epoch: 13 Training Loss: 0.792791 Validation Loss: 0.839745
 Training Accuracy: 63.675000 Validation Accuracy: 60.766667
 Epoch: 14 Training Loss: 0.787382 Validation Loss: 0.826666
 Training Accuracy: 64.014583 Validation Accuracy: 61.275000
 Epoch: 15 Training Loss: 0.786393 Validation Loss: 0.870654
 Training Accuracy: 64.162500 Validation Accuracy: 58.650000
 Epoch: 16 Training Loss: 0.798419 Validation Loss: 0.863961
 Training Accuracy: 63.493750 Validation Accuracy: 60.275000
 Epoch: 17 Training Loss: 0.783020 Validation Loss: 0.878881
 Training Accuracy: 64.277083 Validation Accuracy: 58.891667
 Epoch: 18 Training Loss: 0.781670 Validation Loss: 0.831605
 Training Accuracy: 64.418750 Validation Accuracy: 61.716667
 Epoch: 19 Training Loss: 0.780028 Validation Loss: 0.915492
 Training Accuracy: 64.408333 Validation Accuracy: 55.991667
 Epoch: 20 Training Loss: 0.776635 Validation Loss: 0.826320
 Training Accuracy: 64.731250 Validation Accuracy: 61.233333
 Epoch: 21 Training Loss: 0.796501 Validation Loss: 0.842651
 Training Accuracy: 63.597917 Validation Accuracy: 61.108333
 Epoch: 22 Training Loss: 0.780859 Validation Loss: 0.835370
 Training Accuracy: 64.397917 Validation Accuracy: 61.625000
 Epoch: 23 Training Loss: 0.778767 Validation Loss: 0.820345
 Training Accuracy: 64.720833 Validation Accuracy: 62.083333
 Validation loss decreased (0.824695 --> 0.820345). Saving model ...
 Epoch: 24 Training Loss: 0.779166 Validation Loss: 1.127440
 Training Accuracy: 64.633333 Validation Accuracy: 36.166667
 Epoch: 25 Training Loss: 0.835730 Validation Loss: 0.844096
 Training Accuracy: 60.952083 Validation Accuracy: 60.116667
 Epoch: 26 Training Loss: 0.786725 Validation Loss: 0.831346
 Training Accuracy: 64.439583 Validation Accuracy: 61.166667
 Epoch: 27 Training Loss: 0.777375 Validation Loss: 0.825478
 Training Accuracy: 64.979167 Validation Accuracy: 61.791667
 Epoch: 28 Training Loss: 0.840067 Validation Loss: 0.850965
 Training Accuracy: 61.150000 Validation Accuracy: 60.283333
 Epoch: 29 Training Loss: 0.806351 Validation Loss: 0.876543
 Training Accuracy: 63.060417 Validation Accuracy: 59.841667
 Epoch: 30 Training Loss: 0.787632 Validation Loss: 0.832393

Training Accuracy: 64.231250 Validation Accuracy: 61.600000

FFNN - 66.1% | RNN - 61.4% (Accuracy of models with best validation loss)

RNN accuracy is slightly lower as compared to FFNN. The training loss is decreasing while the validation loss is increasing. Experimented by adding dropout and reducing learning rate but RNN still seems to try to overfit on the data. The discrepancy in accuracy values could be possibly due to RNN only considering 20 words whereas the FFNN averages the entire review.

0.4.2 5(b)

```
[ ]: class GRU(nn.Module):
    def __init__(self, input_size, hidden_size, num_layers, output_size):
        super(GRU, self).__init__()
        self.hidden_size = hidden_size
        self.num_layers = num_layers
        self.gru = nn.GRU(input_size, hidden_size, num_layers,
        ↪batch_first=True, dropout=0.2)
        self.fc = nn.Linear(hidden_size, output_size)

    def forward(self, x):
        batch_size = x.size(0)
        hidden = torch.zeros(self.num_layers, batch_size, self.hidden_size).
        ↪to(device)
        out, hidden = self.gru(x, hidden)
        out = F.log_softmax(self.fc(out[:,-1,:]), dim=1) #Taking the last
        ↪output alone to calculate loss
        return out

gru = GRU(300, 20, 2, 3)
gru = gru.to(device)
loss_fn = nn.CrossEntropyLoss()
optimizer = torch.optim.AdamW(gru.parameters(), lr=0.005)
```

```
[ ]: max_epochs = 30
valid_loss_min = np.Inf

for epoch in range(max_epochs):
    train_loss = 0
    valid_loss = 0
    gru.train()

    train_correct=0
    for batch, labels in train_loader:
        batch, labels = batch.to(device), labels.to(device)
        optimizer.zero_grad()
        output = gru(batch)
        # print(output.shape, '/', labels.shape)
```

```

        loss = loss_fn(output, labels)
        loss.backward()
        optimizer.step()
        train_loss+= loss.item()
        train_correct+= (torch.argmax(output, dim=1)==labels).sum().item()

    gru.eval()
    valid_correct=0
    with torch.no_grad():
        for batch, labels in valid_loader:
            batch, labels = batch.to(device), labels.to(device)
            output = gru(batch)
            loss=loss_fn(output, labels)
            valid_loss+= loss.item()
            valid_correct+= (torch.argmax(output, dim=1)==labels).sum().item()

    train_loss = train_loss/len(train_loader)
    valid_loss = valid_loss/len(valid_loader)
    train_acc = 100 * train_correct / len(train_data)
    valid_acc = 100 * valid_correct / len(valid_data)

    print('Epoch: {} \tTraining Loss: {:.6f} \tValidation Loss: {:.6f}\n
    ↳\tTraining Accuracy: {:.6f} \t Validation Accuracy: {:.6f}'.format(
        epoch+1,
        train_loss,
        valid_loss,
        train_acc,
        valid_acc
    ))

    if valid_loss <= valid_loss_min:
        print('Validation loss decreased ({:.6f} --> {:.6f}). Saving model ...
        ↳'.format(
            valid_loss_min,
            valid_loss))
        torch.save(gru.state_dict(), 'gru.pt')
        valid_loss_min = valid_loss

```

```

Epoch: 1          Training Loss: 0.852240          Validation Loss: 0.782885
Training Accuracy: 59.668750      Validation Accuracy: 63.900000
Validation loss decreased (inf --> 0.782885). Saving model ...
Epoch: 2          Training Loss: 0.758170          Validation Loss: 0.759034
Training Accuracy: 65.779167      Validation Accuracy: 65.200000
Validation loss decreased (0.782885 --> 0.759034). Saving model ...
Epoch: 3          Training Loss: 0.726154          Validation Loss: 0.749555
Training Accuracy: 67.587500      Validation Accuracy: 66.350000
Validation loss decreased (0.759034 --> 0.749555). Saving model ...

```

Epoch: 4 Training Loss: 0.703309 Validation Loss: 0.748379
 Training Accuracy: 68.800000 Validation Accuracy: 66.275000
 Validation loss decreased (0.749555 --> 0.748379). Saving model ...
 Epoch: 5 Training Loss: 0.682284 Validation Loss: 0.742501
 Training Accuracy: 69.950000 Validation Accuracy: 66.800000
 Validation loss decreased (0.748379 --> 0.742501). Saving model ...
 Epoch: 6 Training Loss: 0.665869 Validation Loss: 0.742956
 Training Accuracy: 70.700000 Validation Accuracy: 66.650000
 Epoch: 7 Training Loss: 0.652505 Validation Loss: 0.772581
 Training Accuracy: 71.654167 Validation Accuracy: 65.325000
 Epoch: 8 Training Loss: 0.639380 Validation Loss: 0.760187
 Training Accuracy: 72.245833 Validation Accuracy: 66.441667
 Epoch: 9 Training Loss: 0.625784 Validation Loss: 0.774371
 Training Accuracy: 72.816667 Validation Accuracy: 65.750000
 Epoch: 10 Training Loss: 0.615435 Validation Loss: 0.779434
 Training Accuracy: 73.287500 Validation Accuracy: 66.025000
 Epoch: 11 Training Loss: 0.608425 Validation Loss: 0.765682
 Training Accuracy: 73.835417 Validation Accuracy: 65.916667
 Epoch: 12 Training Loss: 0.597958 Validation Loss: 0.787514
 Training Accuracy: 74.368750 Validation Accuracy: 65.916667
 Epoch: 13 Training Loss: 0.589329 Validation Loss: 0.792497
 Training Accuracy: 74.758333 Validation Accuracy: 65.850000
 Epoch: 14 Training Loss: 0.585013 Validation Loss: 0.811597
 Training Accuracy: 75.120833 Validation Accuracy: 65.950000
 Epoch: 15 Training Loss: 0.576118 Validation Loss: 0.823616
 Training Accuracy: 75.637500 Validation Accuracy: 64.666667
 Epoch: 16 Training Loss: 0.571046 Validation Loss: 0.810080
 Training Accuracy: 75.852083 Validation Accuracy: 65.333333
 Epoch: 17 Training Loss: 0.566572 Validation Loss: 0.841204
 Training Accuracy: 75.881250 Validation Accuracy: 65.541667
 Epoch: 18 Training Loss: 0.558577 Validation Loss: 0.852746
 Training Accuracy: 76.381250 Validation Accuracy: 65.258333
 Epoch: 19 Training Loss: 0.556633 Validation Loss: 0.848621
 Training Accuracy: 76.427083 Validation Accuracy: 65.125000
 Epoch: 20 Training Loss: 0.552553 Validation Loss: 0.839151
 Training Accuracy: 76.752083 Validation Accuracy: 65.091667
 Epoch: 21 Training Loss: 0.544731 Validation Loss: 0.857952
 Training Accuracy: 77.156250 Validation Accuracy: 64.750000
 Epoch: 22 Training Loss: 0.543537 Validation Loss: 0.849694
 Training Accuracy: 77.129167 Validation Accuracy: 63.841667
 Epoch: 23 Training Loss: 0.540733 Validation Loss: 0.864469
 Training Accuracy: 77.168750 Validation Accuracy: 64.833333
 Epoch: 24 Training Loss: 0.534506 Validation Loss: 0.867184
 Training Accuracy: 77.456250 Validation Accuracy: 64.716667
 Epoch: 25 Training Loss: 0.533053 Validation Loss: 0.864193
 Training Accuracy: 77.564583 Validation Accuracy: 64.391667
 Epoch: 26 Training Loss: 0.526711 Validation Loss: 0.874959
 Training Accuracy: 77.825000 Validation Accuracy: 64.550000

Epoch: 27	Training Loss: 0.525434	Validation Loss: 0.901615
Training Accuracy: 78.062500	Validation Accuracy: 64.108333	
Epoch: 28	Training Loss: 0.523132	Validation Loss: 0.892494
Training Accuracy: 78.016667	Validation Accuracy: 63.983333	
Epoch: 29	Training Loss: 0.521680	Validation Loss: 0.907037
Training Accuracy: 78.102083	Validation Accuracy: 63.641667	
Epoch: 30	Training Loss: 0.517928	Validation Loss: 0.903834
Training Accuracy: 78.370833	Validation Accuracy: 64.258333	

GRU - 66.8% (Accuracy after 5 epochs with best validation loss)

Validation loss is increasing with more epochs while training loss is decreasing, showing signs of overfitting. At end of 30 epochs, validation accuracy is 64.25% which is still better than a simple RNN.

0.4.3 5(c)

```
[ ]: class LSTM(nn.Module):
    def __init__(self, input_size, hidden_size, num_layers, output_size):
        super(LSTM, self).__init__()
        self.hidden_size = hidden_size
        self.num_layers = num_layers
        self.lstm = nn.LSTM(input_size, hidden_size, num_layers,
        ↪batch_first=True, dropout=0.2)
        self.fc = nn.Linear(hidden_size, output_size)

    def forward(self, x):
        batch_size = x.size(0)
        hidden = torch.zeros(self.num_layers, batch_size, self.hidden_size).
        ↪to(device)
        cell_state = torch.zeros(self.num_layers, batch_size, self.hidden_size).
        ↪to(device)
        out, (hidden, cell_state) = self.lstm(x, (hidden, cell_state))
        # print(out.shape)
        out = F.log_softmax(self.fc(out[:,-1,:]), dim=1) #Taking the last
        ↪output alone to calculate loss
        return out

lstm = LSTM(300, 20, 2, 3)
lstm = lstm.to(device)
loss_fn = nn.CrossEntropyLoss()
optimizer = torch.optim.AdamW(lstm.parameters(), lr=0.001)
```

```
[ ]: max_epochs = 30
valid_loss_min = np.Inf

for epoch in range(max_epochs):
    train_loss = 0
```

```

valid_loss = 0
lstm.train()

train_correct=0
for batch, labels in train_loader:
    batch, labels = batch.to(device), labels.to(device)
    optimizer.zero_grad()
    output = lstm(batch)
    # print(output.shape, '/', labels.shape)
    loss = loss_fn(output, labels)
    loss.backward()
    optimizer.step()
    train_loss+= loss.item()
    train_correct+= (torch.argmax(output, dim=1)==labels).sum().item()

lstm.eval()
valid_correct=0
with torch.no_grad():
    for batch, labels in valid_loader:
        batch, labels = batch.to(device), labels.to(device)
        output = lstm(batch)
        loss=loss_fn(output, labels)
        valid_loss+= loss.item()
        valid_correct+= (torch.argmax(output, dim=1)==labels).sum().item()

train_loss = train_loss/len(train_loader)
valid_loss = valid_loss/len(valid_loader)
train_acc = 100 * train_correct / len(train_data)
valid_acc = 100 * valid_correct / len(valid_data)

print('Epoch: {} \tTraining Loss: {:.6f} \tValidation Loss: {:.6f}\n
↪\tTraining Accuracy: {:.6f} \t Validation Accuracy: {:.6f}'.format(
    epoch+1,
    train_loss,
    valid_loss,
    train_acc,
    valid_acc
))

if valid_loss <= valid_loss_min:
    print('Validation loss decreased ({:.6f} --> {:.6f}). Saving model ...
↪'.format(
        valid_loss_min,
        valid_loss))
    torch.save(lstm.state_dict(), 'lstm.pt')
    valid_loss_min = valid_loss

```

Epoch: 1 Training Loss: 0.943884 Validation Loss: 0.885581
Training Accuracy: 52.981250 Validation Accuracy: 57.966667
Validation loss decreased (inf --> 0.885581). Saving model ...
Epoch: 2 Training Loss: 0.834866 Validation Loss: 0.825576
Training Accuracy: 61.293750 Validation Accuracy: 61.766667
Validation loss decreased (0.885581 --> 0.825576). Saving model ...
Epoch: 3 Training Loss: 0.797658 Validation Loss: 0.809182
Training Accuracy: 63.581250 Validation Accuracy: 62.716667
Validation loss decreased (0.825576 --> 0.809182). Saving model ...
Epoch: 4 Training Loss: 0.777975 Validation Loss: 0.789377
Training Accuracy: 64.693750 Validation Accuracy: 63.858333
Validation loss decreased (0.809182 --> 0.789377). Saving model ...
Epoch: 5 Training Loss: 0.760450 Validation Loss: 0.779643
Training Accuracy: 65.758333 Validation Accuracy: 64.375000
Validation loss decreased (0.789377 --> 0.779643). Saving model ...
Epoch: 6 Training Loss: 0.745741 Validation Loss: 0.768820
Training Accuracy: 66.468750 Validation Accuracy: 64.991667
Validation loss decreased (0.779643 --> 0.768820). Saving model ...
Epoch: 7 Training Loss: 0.730878 Validation Loss: 0.771671
Training Accuracy: 67.454167 Validation Accuracy: 64.791667
Epoch: 8 Training Loss: 0.721521 Validation Loss: 0.767454
Training Accuracy: 67.639583 Validation Accuracy: 65.216667
Validation loss decreased (0.768820 --> 0.767454). Saving model ...
Epoch: 9 Training Loss: 0.711574 Validation Loss: 0.776195
Training Accuracy: 68.450000 Validation Accuracy: 64.541667
Epoch: 10 Training Loss: 0.701459 Validation Loss: 0.762840
Training Accuracy: 68.931250 Validation Accuracy: 65.566667
Validation loss decreased (0.767454 --> 0.762840). Saving model ...
Epoch: 11 Training Loss: 0.691796 Validation Loss: 0.765836
Training Accuracy: 69.410417 Validation Accuracy: 65.475000
Epoch: 12 Training Loss: 0.685782 Validation Loss: 0.751288
Training Accuracy: 69.700000 Validation Accuracy: 66.100000
Validation loss decreased (0.762840 --> 0.751288). Saving model ...
Epoch: 13 Training Loss: 0.677693 Validation Loss: 0.761019
Training Accuracy: 70.185417 Validation Accuracy: 65.950000
Epoch: 14 Training Loss: 0.671068 Validation Loss: 0.745408
Training Accuracy: 70.512500 Validation Accuracy: 66.375000
Validation loss decreased (0.751288 --> 0.745408). Saving model ...
Epoch: 15 Training Loss: 0.663661 Validation Loss: 0.759947
Training Accuracy: 70.770833 Validation Accuracy: 65.541667
Epoch: 16 Training Loss: 0.658640 Validation Loss: 0.755860
Training Accuracy: 70.827083 Validation Accuracy: 65.866667
Epoch: 17 Training Loss: 0.650487 Validation Loss: 0.764561
Training Accuracy: 71.500000 Validation Accuracy: 65.891667
Epoch: 18 Training Loss: 0.645009 Validation Loss: 0.761992
Training Accuracy: 71.833333 Validation Accuracy: 65.991667
Epoch: 19 Training Loss: 0.639872 Validation Loss: 0.770475
Training Accuracy: 72.125000 Validation Accuracy: 66.391667

Epoch: 20	Training Loss: 0.634310	Validation Loss: 0.763986
Training Accuracy: 72.445833	Validation Accuracy: 66.008333	
Epoch: 21	Training Loss: 0.627621	Validation Loss: 0.780688
Training Accuracy: 72.775000	Validation Accuracy: 65.991667	
Epoch: 22	Training Loss: 0.624477	Validation Loss: 0.776929
Training Accuracy: 72.868750	Validation Accuracy: 66.208333	
Epoch: 23	Training Loss: 0.620246	Validation Loss: 0.820825
Training Accuracy: 73.133333	Validation Accuracy: 64.758333	
Epoch: 24	Training Loss: 0.617020	Validation Loss: 0.778197
Training Accuracy: 73.147917	Validation Accuracy: 65.983333	
Epoch: 25	Training Loss: 0.609854	Validation Loss: 0.776264
Training Accuracy: 73.545833	Validation Accuracy: 65.333333	
Epoch: 26	Training Loss: 0.605385	Validation Loss: 0.793560
Training Accuracy: 73.808333	Validation Accuracy: 66.241667	
Epoch: 27	Training Loss: 0.600010	Validation Loss: 0.784294
Training Accuracy: 74.104167	Validation Accuracy: 66.375000	
Epoch: 28	Training Loss: 0.596917	Validation Loss: 0.782613
Training Accuracy: 74.316667	Validation Accuracy: 66.016667	
Epoch: 29	Training Loss: 0.591855	Validation Loss: 0.787732
Training Accuracy: 74.533333	Validation Accuracy: 66.041667	
Epoch: 30	Training Loss: 0.589975	Validation Loss: 0.788648
Training Accuracy: 74.785417	Validation Accuracy: 66.116667	

RNN - Best model : 61.4% accuracy | GRU - Best model : 66.8% accuracy | LSTM - Best model : 66.37% accuracy

GRU offers significant improvement over RNN due to the mechanism of gates that are capable of learning which inputs are important and need to be remembered. This additional gate mechanism solves the vanishing gradient problem of RNN but is causing the GRU to overfit on training data.

LSTM is computationally slightly more expensive than GRU and offers solutions to the problem of overfitting by adding a forget gate. This results in better performance in both training and validation data.

0.5 References

1. https://radimrehurek.com/gensim/auto_examples/tutorials/run_word2vec.html
2. <https://www.kaggle.com/code/mishra1993/pytorch-multi-layer-perceptron-mnist/notebook>
3. <https://pytorch.org/docs/stable/generated/torch.nn.RNN.html>
4. <https://pytorch.org/docs/stable/generated/torch.nn.GRU.html>
5. <https://pytorch.org/docs/stable/generated/torch.nn.LSTM.html>